

PyMuPDF Documentation

version 1.8

Ruikai Liu

Jorj McKie

March 07, 2016

Contents

The PyMuPDF Documentation	1
Introduction	1
Note on the Name <code>fitz</code>	1
Installation	2
Step 1: Download PyMuPDF	2
Step 2: Download and Generate MuPDF 1.8	2
Step 3: Build / Setup PyMuPDF	2
Note on using UPX	2
Tutorial	3
Import the Bindings	3
Open a Document	3
Some Document methods and attributes	3
Access Meta Data	3
Work with Outlines	4
Work with Pages	4
Inspect the links on a Page	4
Render a Page	4
Save the page image in a file	4
Display the image in dialog managers	5
Text extraction	5
Text Searching	5
Output	5
Close	6
Example: Dynamically cleaning up corrupt PDF documents	7
Classes	8
Colorspace	8
Device	9
DisplayList	10
Document	11
Identity	15
IRect	16
Link	18
linkDest	19
Matrix	21
Shifting	23
Flipping	23
Shearing	24
Rotating	25
Outline	26
Page	27

Pixmap	29
Details on Saving Images with <code>writeImage()</code>	32
Pixmap Example Code Snippets	32
Point	35
Rect	36
TextPage	38
TextSheet	39
Constants and Enumerations	40
Constants	40
Enumerations	40
Appendix 1: Performance	42
Part 1: Parsing	42
Part 2: Text Extraction	45
Part 3: Image Rendering	46
Appendix 2: Details on Text Extraction	47
General structure of a TextPage	47
Output of <code>getText(output="text")</code>	47
Output of <code>getText(output="html")</code>	47
Output of <code>getText(output="json")</code>	47
Output of <code>getText(output="xml")</code>	48
Resource Requirements	49
Performance	49
Data Sizes	49
Index	51

The PyMuPDF Documentation

Introduction

PyMuPDF (formerly known as **python-fitz**) is a Python binding for **MuPDF** - "a lightweight PDF and XPS viewer".

MuPDF can access files in PDF, XPS, OpenXPS, CBZ (comic book) and EPUB (e-book) formats.

These are files with extensions `*.pdf`, `*.xps`, `*.oxps`, `*.cbz` or `*.epub` (so in essence, with this binding you can develop **e-book viewers in Python** ...)

PyMuPDF provides access to all important functions of MuPDF from within a Python environment. Nevertheless, we are continuously expanding this function set.

MuPDF stands out among all similar products for its top rendering capability and unsurpassed processing speed.

You can check this out yourself: Compare the various free PDF-viewers. In terms of speed and rendering quality **SumatraPDF** ranges at the top (apart from MuPDF's own standalone viewer) - and it is based on MuPDF!

While PyMuPDF has been available since several years for an earlier version of MuPDF (1.2), it was until only mid May 2015, that its creator and a few co-workers decided to elevate it to support the current release of MuPDF (first 1.7a and, since November 2015, 1.8).

And we are determined to keep PyMuPDF current with future MuPDF changes!

This work is now completed.

PyMuPDF has been tested on Linux, Windows 7, Windows 10, Python 2 and Python 3 (x86 versions). Other platforms should work too as long as MuPDF supports them.

The main differences compared to version 1.2 are

- A greatly simplified installation procedure: For Windows and Linux platforms it should come down to running the `python setup.py install` command.
- The API has changed: it is now simpler and a lot less cryptic.
- The supported function set has been significantly increased: apart from rendering, MuPDF's traditional strength, we now also offer a wide range of text extraction options.
- Demo code has been extended, and an additional `examples` directory is there to contain working programs. Among them are an editor for a document's table of contents, a full featured document joiner and a document-to-text conversion utility.

We invite you to join our efforts by contributing to the the wiki pages, by using what is there - and, of course, by submitting issues and bugs to the site!

Note on the Name `fitz`

The Python import statement for this library is `import fitz`. Here is the reason why:

The original rendering library for MuPDF was called `Libart`. "After Artifex Software acquired the MuPDF project, the development focus shifted on writing a new modern graphics library called `Fitz`. Fitz was originally intended as an R&D project to replace the aging Ghostscript graphics library, but has instead become the rendering engine powering MuPDF." (Quoted from [Wikipedia](#)).

Installation

This describes how to install PyMuPDF.

Step 1: Download PyMuPDF

Download this repository and unzip / decompress it. This will give you a folder, let us call it `PyFitz`.

Step 2: Download and Generate MuPDF 1.8

If your platform is not Windows (or if it is and you want to generate MuPDF) download `mupdf-1.8-source.tar.gz` from [MuPDF version 1.8 source](#), now and unzip / decompress it. Let us call the resulting folder `mupdf`.

Put it inside `PyFitz` as a subdirectory, if you want to keep everything in one place.

If your platform is **not Windows**, you must **generate MuPDF now**. The MuPDF download includes generation procedures / makefiles for numerous platforms.

On Windows, you have two options:

- if you **have installed MS Visual Studio**, generate the MuPDF `lib` files now. The respective VS project files are contained in `./PyFitz/mupdf/platform/win32`. If that worked fine, the `lib` files are now in `./PyFitz/mupdf/platform/win32/Release`. Update `setup.py` to reference this directory as `library_dirs=['./mupdf/platform/win32/Release']`. Please note that your VS version needs to use a C++ compiler compatible with the compiler that was used to generate your Python version. Generally, this is MSVC 9.0 (contained in VS2008 and Microsoft Visual C++ for Python) for Python 2.7, or MSVC 10.0 (contained in VS2010, VS Community 2013 and Microsoft Windows SDK for Windows 7 and .NET Framework 4 (ISO)) for Python up to 3.4.
- if you have not installed Visual Studio or if you do not want to generate MuPDF, you must download [PyMuPDF Optional Material](#) now and unzip / decompress its content in directory `./PyFitz/PyMuPDF-optional-material`. This optional material contains the `lib` files needed for generating PyMuPDF for your Python, and the MuPDF header files. Update your `setupxy.py` version, parameter `include_dirs`, to point to these header files.

Supporting Python 3.5 on Windows is still work in progress.

Step 3: Build / Setup PyMuPDF

If necessary, adjust the `setup.py` script now. E.g. make sure that

- the include directory is correctly set in sync with your directory structure
- the object code libraries are correctly defined

Now perform a `python setup.py install`

Note on using UPX

In Windows systems, your PyMuPDF installation will end up with three files: `__init__.py`, `fitz.py` and `_fitz.pyd` in the `site-packages` directory (PYD files are Python's DLL version on Windows systems). `_fitz.pyd` has a size of 9.5 to 10 MB.

You can reduce this by applying the compression utility UPX to it: `upx -9 _fitz.pyd`. This will reduce the file size to about 4.5 MB. This should reduce load times (`import fitz` statement) while keeping it fully functional.

Tutorial

This tutorial will show you the use of MuPDF in Python step by step.

Because MuPDF supports not only PDF, but also XPS, OpenXPS and EPUB formats, so does PyMuPDF. Nevertheless we will only talk about PDF files for the sake of brevity.

As for string handling, MuPDF will pass back any string as UTF-8 encoded - no exceptions. Where this binding has added functionality, we usually decode string to unicode. An example is the `Document.ToC()` method.

Import the Bindings

The Python bindings to MuPDF are made available by this import statement:

```
import fitz
```

Open a Document

In order to access a supported document, it must be opened with the following statement:

```
doc = fitz.Document(filename)
```

This will create `doc` as a `Document` object. `filename` must be a Python string or unicode object that specifies the name of an existing file (with or without a fully or partially qualified path).

It is also possible to construct a document from memory data, i.e. without using a file. See `Document` for details.

A `Document` contains several attributes and functions. Among them are meta information (like "author" or "subject"), number of total pages, outline and encryption information.

Some Document methods and attributes

Method / Attribute	Description
<code>Document.pageCount</code>	Number of pages of filename (integer).
<code>Document.metadata</code>	Metadata of the Document (dictionary).
<code>Document.outline</code>	First outline entry of Document
<code>Document.getToC()</code>	Table of contents of Document (list).
<code>Document.loadPage()</code>	Create a <code>Page</code> object.

Access Meta Data

`Document.metadata` is a Python dictionary with the following keys. For details of their meanings and formats consult the PDF manuals, e.g. [Adobe PDF Reference sixth edition 1.7 November 2006](#). Further information can also be found in chapter `Document`. The meta data fields are of type string if not otherwise indicated and may be missing, in which case they contain `None`.

Key	Value
<code>producer</code>	Producer (producing software)
<code>format</code>	PDF format, e.g. 'PDF-1.4'
<code>encryption</code>	Encryption method used
<code>author</code>	Author
<code>modDate</code>	Date of last modification

keywords	Keywords
title	Title
creationDate	Date of creation
creator	Creating application
subject	Subject

Work with Outlines

The easiest way to get all outlines of a document, is creating a table of contents:

```
toc = doc.getToC(simple = True) # the simple form, if False, link information is inclu
```

This will return a Python list `[[level, title, page, link], ...]` (or `[]`).

`level` is the hierarchy level of the entry (starting from 1), `title` is the entry's title (unicode), and `page` the page number (1-based). `link` is present if `simple = False` is specified. Its meaning can be look up under `Page.getLinks()`.

If you want a more detailed control of what you get, enter the document's outline tree like this:

```
olItem = doc.outline # the document's first outline item
```

This creates `olItem` as an [Outline](#) object. Look there for further details.

Work with Pages

Tasks that can be performed with a [Page](#) are at the core of MuPDF's functionality. Among other things, you can render a [Page](#), optionally zooming, rotating or shearing it. You can write it's image to files (in PNG format), extract text from it or search for text strings. At first, a page object must be created:

```
page = doc.loadPage(n) # represents page n of the document
```

Some typical uses of [Page](#) objects follow:

Inspect the links on a Page

Here is how to get all links and their types:

```
#-----
# Get all links of the current page
#-----
links = page.getLinks()
```

`links` is a Python list containing Python dictionaries as entries. For details see `Page.getLinks()`.

Render a Page

This example creates an image out of a page's content:

```
pix = page.getPixmap(matrix = fitz.Identity, colorspace = "RGB")
#-----
# now pix contains an RGB image of the page, ready to be used
#-----
```

Save the page image in a file

We can simply store the image in a PNG file:


```
pix.writePNG("test.png")
```

Display the image in dialog managers

We can also use the image in a dialog. `Pixmap.samples` represents the area of bytes of all the pixels as a Python bytearray. This area (or its `str()`-version), is directly usable by presumably most dialog managers. Here are two examples. Please also have a look at the examples directory of this repository.

wxPython:

```
data = pix.samples                                # data = bytearray of raw pixel data (RGBA)
bitmap = wx.BitmapFromBufferRGBA(pix.width, # image width
                                pix.height, # image height
                                str(data))    # wxPython only accepts strings, no bytearrays
```

Tkinter:

```
data = pix.samples
# the following requires: "from PIL import Image"
img = Image.frombytes("RGBA", [pix.width, pix.height], str(data))
photo = ImageTk.PhotoImage(img)
```

Text extraction

We can also extract all text of a page in one chunk of string:

```
text = page.getText(output = "text")
```

For the `output` parameter, the following values can be specified:

- `text`: plain text with line breaks. No format and no position info.
- `html`: line breaks, alignment, grouping. No format and no position info
- `json`: full formatting info (except colors and fonts) down to spans (see appendix)
- `xml`: full formatting info (except colors) down to individual characters

To give you an idea about the output of these alternatives, we did text extracts from this document's PDF version and several other examples. See the appendix for details about implications on processing times and space requirements.

Text Searching

You can find out, exactly where on a page a certain string appears like this:

```
areas = page.searchFor("mupdf", hit_max = 32)
```

The variable `areas` will now contain a list of up to 32 `Rect` rectangles which surround each occurrence of string "mupdf" (not case sensitive).

Please also do have a look at the demonstration program `demo.py`. Among others it contains details on how the `TextPage`, `TextSheet`, `Device` and `DisplayList` classes can be used for a more direct control, e.g. when performance considerations require it.

Output

Output capabilities of MuPDF (such as PDF generation) are currently very limited. However, a copy of a currently opened document can be created (PDF only, none of the other file types).

We support this with the method `Document.save()`.

If the document had been successfully decrypted before, `save()` will automatically create a decrypted copy.

In addition, this method will also perform some clean-up:

If the document contains invalid or broken xrefs, the saved version will have them automatically corrected, which makes it readable by other Python PDF software, like [pdfcrowd](#) or [PyPDF2](#). In many cases, the saved version will also be smaller than the original.

`Document.save()` now also supports all options of MuPDF's command line utility `mutool clean`, see below (`mutool clean` option = "MCO").

Option	MCO	Effect
<code>garbage = 1</code>	<code>-g</code>	garbage collect unused objects
<code>garbage = 2</code>	<code>-gg</code>	in addition to 1, compact xref tables
<code>garbage = 3</code>	<code>-ggg</code>	in addition to 2, merge duplicate objects
<code>clean = 1</code>	<code>-s</code>	clean content streams (avoid or use with care!)
<code>deflate = 1</code>	<code>-z</code>	deflate uncompressed streams
<code>ascii = 1</code>	<code>-a</code>	convert data to ASCII format
<code>linear = 1</code>	<code>-l</code>	create a linearized version (do not use yet)
<code>expand = 1</code>	<code>-i</code>	decompress images
<code>expand = 2</code>	<code>-f</code>	decompress fonts
<code>expand = 255</code>	<code>-d</code>	decompress all
<code>incremental = 1</code>	n/a	only save data that have changed (do not use yet)

Please note, that `Document.save()`, according to MuPDF's documentation, is still being further developed, so expect changes in the future here. Also observe our remarks in the above table.

Like with `mutool clean`, not all combinations of the above options may work for all documents - so be ready to experiment a little.

We have found, that the fastest, yet very stable combination is `mutool clean -ggg -z`, giving good compression results. In PyMuPDF this corresponds to `doc.save(filename, garbage=3, deflate=1)`.

In some cases, best compression factors result, if `expand` and `deflate` are used together, though they seem to be contradictory. This works, because MuPDF is forced to expand and then re-compress all objects, which will correct poor compressions during document creation.

Close

In some situations it is desirable to "close" a [Document](#) such that it becomes fully available again to the OS while your program is still running.

This can be achieved by the `Document.close()` method. Apart from closing the file, all buffer areas associated with the document will be freed. If the document has been created from memory data, no underlying file is opened by MuPDF, so only the buffer release will take place.

Caution:

As with normal file objects, after close, the document and all objects referencing it will be invalid and **must no longer be used**. This binding protects against most such invalid uses by disabling properties and methods of the [Document](#) and any associated `Document.loadPage()` objects.

However, re-opening a previously closed file by a new [Document](#) is no problem. Please also do have a look at the following valid example:

```
doc = fitz.Document(f_old)           # open a document
<... some statements ...>           # e.g. decryption
doc.save(fnew, garbage=3, deflate=1) # save a decrypted / compressed version
doc.close()                         # close input file
os.remove(f_old)                    # remove it
os.rename(f_new, f_old)              # rename the decrypted / cleaned version
doc = fitz.Document(f_old)           # use it as input for MuPDF
```

Example: Dynamically cleaning up corrupt PDF documents

This shows a potential use of PyMuPDF with another Python PDF library (pdfwr).

If a PDF is broken or needs to be decrypted, one could dynamically invoke PyMuPDF to recover from problems like so:

```
import sys
from pdfwr import PdfReader
import fitz
from cStringIO import StringIO

#-----
# 'tolerant' PDF reader
#-----
def reader(fname):
    ifile = open(fname, "rb")
    idata = ifile.read() # put in memory
    ifile.close()
    ibuffer = StringIO(idata) # convert to stream
    try:
        return PdfReader(ibuffer) # let us try
    except: # problem! see if PyMuPDF can heal it
        doc = fitz.Document("application/pdf",
                             idata,
                             len(idata)) # scan pdf data in memory
        doc.save("test.pdf", # may want to use a temp file
                garbage=3,
                deflate=1) # save a cleaned version
        ifile = open("test.pdf", "rb") # open it
        idata = ifile.read() # put in memory
        ifile.close()
        ibuffer = StringIO(idata) # convert to stream
        return PdfReader(ibuffer) # now let pdfwr retry

#-----

pdf = reader(sys.argv[1])
print pdf.Info
# do further processing
```

With the command line utility `pdftk` a similar result can be achieved, see [here](#). It even supports buffers for input **and** output. However you must invoke it as a separate process via `subprocess.Popen`, using `stdin` and `stdout` as communication vehicles.

Classes

Colorspace

Represents the color space of a [Pixmap](#).

Class API

class Colorspace

`__init__` (self, colorspace)
Constructor

Parameters: **colorspace** (*int*) -- A number identifying the colorspace. Possible values are `CS_RGB`, `CS_GRAY` and `CS_CMYK`.

Device

The different format handlers (pdf, xps, etc.) interpret pages to a "device". These devices are the basis for everything that can be done with a page: rendering, text extraction and searching. The device type is determined by the selected construction method.

Class API

class Device

`__init__` (self, object)

Constructor for either a pixel map or a display list device.

object

An object representing one of [Pixmap](#), or [DisplayList](#)

Type: instance

`__init__` (self, textsheet, textpage)

Constructor for a text page device.

textsheet

A [TextSheet](#) object.

Type: instance

textpage

A [TextPage](#) object.

Type: instance

DisplayList

DisplayList is a list containing drawing commands (text, images, etc.). The intent is two-fold:

1. as a caching-mechanism to reduce parsing of a page
2. as a data structure in multi-threading setups, where one thread parses the page and another one renders pages.

A `DisplayList` is populated with objects from a page by running `Page.run()` on a `Device`. Replay the list (once or many times) by invoking the display list's `run()` function.

Method	Short Description
<code>run()</code>	(Re)-run a display list through a device.

Class API

`class DisplayList`

`fitz.DisplayList (self)`

Create a rendering device for a display list.

When the device is rendering a page it will populate the display list with drawing commands (text, images, etc.). The display list can later be reused to render a page many times without having to re-interpret the page from the document file.

Return type: `Device`

`run (self, dev, ctm, area)`

Parameters:

- **dev** (`Device`) -- Device obtained from `Device`
- **ctm** (`Matrix`) -- Transform matrix to apply to display list contents.
- **area** (`Rect`) -- Only the part of the contents of the display list visible within this area will be considered when the list is run through the device. This does not imply for tile objects contained in the display list.

Document

This class represents a document. It can be constructed from a file or from memory. See below for details.

Method / Attribute	Short Description
<code>Document.authenticate()</code>	decrypts the document
<code>Document.loadPage()</code>	reads a page
<code>Document.save()</code>	saves a copy of the document
<code>Document.getToC()</code>	create a table of contents
<code>Document.getPagePixmap()</code>	create a pixmap from a page by its number
<code>Document.getPageText()</code>	extracts the text of a page by its number
<code>Document.close()</code>	closes the document
<code>Document.isClosed</code>	has document been closed?
<code>Document.outline</code>	first <i>Outline</i> item
<code>Document.name</code>	filename of document
<code>Document.needsPass</code>	require password to access data?
<code>Document.isEncrypted</code>	is document still encrypted?
<code>Document.pageCount</code>	the document's number of pages
<code>Document.metadata</code>	the document's meta data

Class API

class Document

`__init__(self, filename)`

Constructs a `Document` object from a file.

Parameters: `filename` (*string*) -- A string (UTF-8 or unicode) containing the path / name of the document file to be used. The file will be opened and remain open until either explicitly closed (see below) or until end of program.

Return type: `Document`

Returns: A `Document` object.

`__init__(self, filetype, stream=data, streamlen=len(data))`

Constructs a `Document` object from memory data.

Parameters:

- **filetype** (*string*) -- A string specifying the type of document contained in `stream`. This may be either something that looks like a filename (e.g. `x.pdf`), in which case MuPDF uses the extension to determine the type, or a mime type like `application/pdf`. Recommended is using the filename scheme, or even the name of the original file for documentation purposes.
- **stream** (*string*) -- A string of data representing the content of a supported document type.
- **streamlen** (*int*) -- An integer specifying the length of the stream.

Return type: `Document`

Returns: A `Document` object.

`authenticate(password)`

Decrypts the document with the string `password`. If successful, all of the document's data can be accessed (e.g. for rendering).

Parameters: `password` (*string*) -- The password to be used.

Return type: `int`

Returns: `True (1)` if decryption with `password` was successful, `False (0)` otherwise.

loadPage (`number`)

Loads a `Page` for further processing like rendering, text searching, etc. See the [Page](#) object.

Parameters: `number (int)` -- page number, zero-based (0 is the first page of the document).

Return type: [Page](#)

getToC (`simple = True`)

Creates a table of contents out of the document's outline chain.

Parameters: `simple (boolean)` -- Indicates whether a detailed ToC is required. If `simple = False` is specified, each entry of the list also contains a dictionary with [Link](#) details for each outline entry.

Return type: `list`

getPagePixmap (`pno, matrix = fitz.Identity, colorspace = "rgb"`)

Creates a pixmap from one of the document's pages - identified by number `pno` (zero-based).

Parameters:

- `pno (int)` -- Page number, zero-based
- `matrix (Matrix)` -- A transformation matrix - default is [Identity](#).
- `colorspace (string)` -- A string specifying the requested colorspace - default is `rgb`.

Return type: [Pixmap](#)

getPageText (`pno, output = "text"`)

Extracts the text of a page given its page number `pno` (zero-based).

Parameters:

- `pno (int)` -- Page number, zero-based
- `output (string)` -- A string specifying the requested output format: `text`, `html`, `json` or `xml`. Default is `text`.

Return type: `String`

save (`outfile, garbage=0, clean=0, deflate=0, incremental=0, ascii=0, expand=0, linear=0`)

Saves a copy of the document under the name `outfile`. Include path specifications as necessary. Only PDF documents are supported by this function. Internally the document may have changed. E.g. after a successful `authenticate`, a decrypted copy will be saved, and, in addition (even without any of the optional parameters), some basic cleaning of the document data could also have occurred, e.g. broken xref tables have been corrected and incremental changes have been resolved.

Parameters:

- `outfile (string)` -- The file name to save to. Must be different from the original `filename / filetype` value or else a `ValueError` will be raised.
- `garbage (int)` -- Do garbage collection: 0 = none, 1 = remove unused objects, 2 = in addition compact xref tables, 3 = in addition merge duplicate objects.
- `clean (int)` -- Clean content streams: 0 = False, 1 = True.
- `deflate (int)` -- Deflate uncompressed streams: 0 = False, 1 = True.
- `incremental (int)` -- Only save changed objects: 0 = False, 1 = True.
- `ascii (int)` -- Where possible make the output ASCII: 0 = False, 1 = True.
- `expand (int)` -- One byte bitfield to decompress contents: 0 = none, 1 = images, 2 = fonts, 255 = all. This convenience option generates a decompressed file version that can be better read by some other programs.
- `linear (int)` -- Save a linearised version of the document: 0 = False, 1 = True. This option creates a file format for improved performance when read via internet connections.

Return type: int

Returns: Count of errors that occurred during save. Note: PyMuPDF will recover from many errors encountered in a PDF and continue processing.

close ()

Releases space allocations associated with the document, and, if created from a file, closes `filename` thus releasing control of it to the OS.

outline

Contains either `None` or the first `Outline` entry of the document. Can be used as a starting point to walk through all outline items. If an encrypted document has not yet been authenticated, an `AttributeError` exception will be raised, when this attribute is being accessed.

Return type: `Outline`

isClosed

`False` (0) if document is still open, `True` (1) otherwise. If closed, most other attributes and all methods will have been deleted / disabled. In addition, `Page` objects referring to this document (i.e. created with `Document.loadPage()`) will no longer be usable. For reference purposes, `Document.name` still exists and will contain the filename of the original document.

Return type: int

needsPass

Contains an indicator showing whether the document is encrypted (`True` = 1) or not (`False` = 0). This indicator remains unchanged - even after the document has been authenticated.

Return type: bool

isEncrypted

This indicator initially equals the value of `needsPass`. After a successful authentication, it is set to `False` = 0 to reflect the situation.

Return type: bool

metadata

Contains the document's meta data as a Python dictionary or `None` if the document is encrypted. Its keys are `format`, `encryption`, `title`, `author`, `subject`, `keywords`, `creator`, `producer`, `creationDate`, `modDate`. All item values are strings or `None`.

Except `format` and `encryption`, the key names correspond in an obvious way to a PDF's "official" meta data fields `/Creator`, `/Producer`, `/CreationDate`, `/ModDate`, `/Title`, `/Author`, `/Subject`, `/Keywords` respectively.

The value of `format` contains the version of the PDF format (e.g. 'PDF-1.6').

The value of `encryption` either contains `None` (not encrypted), or a string naming the used encryption method (e.g. 'Standard V4 R4 128-bit RC4').

If the date fields contain meaningful data (which need not be the case), they are strings in the PDF-internal timestamp format "D:<TS><TZ>", where

<TS> is the 12 character ISO timestamp `YYYYMMDDhhmmss` (`YYYY` - year, `MM` - month, `DD` - day, `hh` - hour, `mm` - minute, `ss` - second), and

<TZ> is a time zone value (time interval relative to GMT) containing a sign ('+' or '-'), the hour (`hh`), and the minute ('`mm`', attention: enclose in apostrophies!).

For example, a Venezuelan value might look like `D:20150415131602-04'30'`, which corresponds to the timestamp April 15, 2015, at 1:16:02 pm local time Venezuela.

Return type: dict

name

Contains the `filename` or `filetype` value with which `Document` was created.

Return type: string

pageCount

Classes

Contains the number of pages of the document. May return 0 for documents with no pages.

Return type: int

Identity

Identity is just a [Matrix](#) that performs no action, to be used whenever the syntax requires a [Matrix](#), but no actual transformation should take place.

Caution: `Identity` is a constant in the C code and therefore **readonly, do not try to modify** its properties in any way, i.e. you must not manipulate its `[a,b,c,d,e,f]`, neither apply any method.

`Matrix(1, 1)` creates a matrix that acts like `Identity`, but it may be changed. Use this when you need a starting point for further modification, e.g. by one of the [Matrix](#) methods.

In other words:

```
# the following will not work - the interpreter will crash!
m = fitz.Identity.preRotate(90)

# do this instead:
m = fitz.Matrix(1, 1).preRotate(90)
```

IRect

IRect is a rectangular bounding box similar to [Rect](#), except that all corner coordinates are integers. IRect is used to specify an area of pixels, e.g. to receive image data during rendering.

Attribute / Method	Short Description
<code>IRect.getRect()</code>	return arectangle with same coordinates
<code>IRect.width</code>	Width of the bounding box
<code>IRect.height</code>	Height of the bounding box
<code>IRect.x0</code>	X-coordinate of the top left corner
<code>IRect.y0</code>	Y-coordinate of the top left corner
<code>IRect.x1</code>	X-coordinate of the bottom right corner
<code>IRect.y1</code>	Y-coordinate of the bottom right corner

Class API

class IRect

`__init__ (self, x0, y0, x1, y1)`

Constructor. Without parameters defaulting to `IRect(0, 0, 0, 0)`, an empty rectangle. Also see the example below. Function `Rect.round()` creates the smallest IRect containing Rect.

`getRect ()`

A convenience function returning a [Rect](#) with the same coordinates as floating point values.

`width`

Contains the width of the bounding box. Equals $x1 - x0$.

Type: int

`height`

Contains the height of the bounding box. Equals $y1 - y0$.

Type: int

`x0`

X-coordinate of the top left corner.

Type: int

`y0`

Y-coordinate of the top left corner.

Type: int

`x1`

X-coordinate of the bottom right corner.

Type: int

`y1`

Y-coordinate of the bottom right corner.

Type: int

Example:

```
Python 2.7.9 (default, Dec 10 2014, 12:24:55) [MSC v.1500 32 bit (Intel)] on win32
Type "copyright", "credits" or "license()" for more information.
>>> import fitz
```

Classes

```
>>> irect = fitz.IRect(10, 10, 410, 610)
>>> irect
fitz.IRect(10, 10, 410, 610)
>>> irect.height
600
>>> irect.width
400
>>>
```

Link

Represents a pointer to somewhere (this document, other documents, the internet). Links exist per document page, and they are forward-chained to each other, starting from an initial link which is accessible by the `Page.loadLinks()` method.

Attribute	Short Description
<code>Link.rect</code>	Clickable area in untransformed coordinates.
<code>Link.dest</code>	Kind of link destination.
<code>Link.next</code>	Link to next link

Class API

`class Link`

rect

The area that can be clicked in untransformed coordinates.

Return type: [Rect](#)

dest

The link destination kind. An integer to be interpreted as one of the `FZ_LINK_*` values.

Return type: `int`

next

The next `Link` or `None`

Return type: [Link](#)

linkDest

Class representing the *dest* property of an outline entry.

Attribute	Short Description
<code>linkDest.dest</code>	Destination
<code>linkDest.fileSpec</code>	File specification (path, filename)
<code>linkDest.flags</code>	Descriptive flags
<code>linkDest.isMap</code>	Is this a MAP?
<code>linkDest.isUri</code>	Is this an URI?
<code>linkDest.kind</code>	Kind of destination
<code>linkDest.lt</code>	Top left coordinates
<code>linkDest.named</code>	Name if named destination
<code>linkDest.newWindow</code>	Name of new window
<code>linkDest.page</code>	Page number
<code>linkDest.rb</code>	Bottom right coordinates
<code>linkDest.uri</code>	URI

Class API

class linkDest

dest

Destination of linkDest.

Return type: [Link](#)

fileSpec

Contains the filename (including any path specifications) this link points to, if applicable.

Return type: `string`

flags

A one-byte bitfield consisting of indicators describing the validity and meaning of the different aspects of the destination. As far as possible, link destinations are constructed such that e.g. `linkDest.lt` and `linkDest.rb` can be treated as defining a bounding box, though the validity flags (see `LINK_FLAG_*` values) indicate which of the values were actually specified. Note that the numerical values for each of the `LINK_FLAGS` are powers of 2 and thus indicate the position of the bit to be tested. More than one bit can be `True`, so do not test for the value of the integer.

Return type: `int`

isMap

This flag specifies whether to track the mouse position when the URI is resolved. Default value: `False`.

Return type: `bool`

isUri

Specifies whether this destination is an internet resource.

Return type: `bool`

kind

Indicates the type of this destination, like a place in this document, a URI, a file launch, an action or a place in another file. Look at index entries `FZ_LINK_*` to see the names and numerical values.

Return type: `int`

lt

The top left [Point](#) of the destination.

Return type: [Point](#)

named

This destination refers to some named resource of the document (see Adobe PDF documentation).

Return type: `int`

newWindow

This destination refers to an action that will open a new window.

Return type: `bool`

page

The page number (in this or the target document) this destination points to.

Return type: `int`

rb

The bottom right [Point](#) of this destination.

Return type: [Point](#)

uri

The name of the URI this destination points to.

Return type: `string`

Matrix

Matrix is a row-major 3x3 matrix used by image transformations in MuPDF. With matrices you can manipulate the rendered image of a page in a variety of ways: (parts of) the page can be rotated, zoomed, flipped, sheared and shifted by setting some or all of just six numerical values.

Since all points or pixels live in a two-dimensional space, one column vector of that matrix is a constant unit vector, and only the remaining six elements are used for manipulations. These six elements are usually represented by $[a, b, c, d, e, f]$. Here is how they are positioned in the matrix:

$$\begin{bmatrix} a & b & 0 \\ c & d & 0 \\ e & f & 1 \end{bmatrix}$$

It should be noted, that

- the below methods are just convenience functions - everything they do, can also be achieved by directly manipulating $[a, b, c, d, e, f]$
- all manipulations can be combined - you can construct a matrix that does a rotate **and** a shear **and** a scale **and** a shift etc. in one go

Method / Attribute	Description
<code>Matrix.__init__()</code>	Constructor.
<code>Matrix.preRotate()</code>	Perform a rotation
<code>Matrix.preScale()</code>	Perform a scaling
<code>Matrix.preShear()</code>	Perform a shearing
<code>Matrix.a</code>	Zoom factor X direction
<code>Matrix.b</code>	Shearing effect Y direction
<code>Matrix.c</code>	Shearing effect X direction
<code>Matrix.d</code>	Zoom factor Y direction
<code>Matrix.e</code>	Horizontal shift
<code>Matrix.f</code>	Vertical shift

Class API

`class Matrix`

`__init__(self, a=1, b=0, c=0, d=1, e=0, f=0)`

Constructor. `Matrix(1, 1)` will construct a modifiable version of the [Identity](#) matrix.

`preRotate(deg)`

Modify the matrix to perform a clockwise rotation for positive `deg` degrees, else counterclockwise. The matrix elements will change in the following way:

$[a, b, c, d, e, f] \rightarrow [\cos(\deg), \sin(\deg), -\sin(\deg), \cos(\deg), e, f]$.

Parameters: `deg (float)` -- The rotation angle in degrees (use conventional notation based on $\pi = 180$ degrees).

`preScale(sx, sy)`

Modify the matrix to scale by the zoom factors `sx` and `sy`. Has effects on attributes `a` and `d` only.

Parameters:

- `sx (float)` -- Zoom factor in X direction. For the effect see description of attribute `a`.
- `sy (float)` -- Zoom factor in Y direction. For the effect see description of attribute `d`.

preShear (sx, sy)

Modify the matrix to perform a shearing, i.e. transformation of rectangles into parallelograms (rhomboids). Has effects on attributes `b` and `c` only.

Parameters:

- **sx** (*float*) -- Shearing effect in X direction. See attribute `c`.
- **sy** (*float*) -- Shearing effect in Y direction. See attribute `b`.

- a** Scaling in X-direction (**width**). For example, a value of 0.5 performs a shrink of the **width** by a factor of 2. If `a < 0`, a left-right flip will (additionally) occur.

Type: float

- b** Causes a shearing effect: each `Point(x, y)` will become `Point(x, y - b*x)`. Therefore, looking from left to right, e.g. horizontal lines will be "tilt" - downwards if `b > 0`, upwards otherwise (`b` is the tangens of the tilting angle).

Type: float

- c** Causes a shearing effect: each `Point(x, y)` will become `Point(x - c*y, y)`. Therefore, looking upwards, vertical lines will be "tilt" - to the left if `c > 0`, to the right otherwise (`c` is the tangens of the tilting angle).

Type: float

- d** Scaling in Y-direction (**height**). For example, a value of 1.5 performs a stretch of the **height** by 50%. If `d < 0`, an up-down flip will (additionally) occur.

Type: float

- e** Causes a horizontal shift effect: Each `Point(x, y)` will become `Point(x + e, y)`. Positive (negative) values of `e` will shift right (left).

Type: float

- f** Causes a vertical shift effect: Each `Point(x, y)` will become `Point(x, y - f)`. Positive (negative) values of `f` will shift down (up).

Type: float

Examples

Here are examples to illustrate some of the effects achievable with matrices. The following pictures start with a page of the PDF version of this help file. We show what will happen when a matrix is being applied (though always full pages are created, only parts are displayed here to save space).

This is the original page image

Classes

Matrix

Matrix is a row-major 3x3 matrix used for representing transformations of coordinates throughout MuPDF.

Since all points or pixels reside in a two-dimensional space, one column vector of the matrix is the constant unit vector, and only the remaining six elements may vary. These six elements are usually represented by $[a, b, c, d, e, f]$. Here is how they are positioned in the matrix:

$$\begin{bmatrix} a & b & 0 \\ c & d & 0 \\ e & f & 1 \end{bmatrix}$$

It should be noted, that the below methods are just convenience functions. Each of them manipulates some of the six matrix elements in a specific way. By directly changing $[a, b, c, d, e, f]$, any of these functions can be replaced.

Shifting

We transform it with a matrix where $e = 100$ (right shift by 100 pixels)

Classes

Matrix is a row-major 3x3 matrix used for representing transformations of coordinates throughout MuPDF.

Since all points or pixels reside in a two-dimensional space, one column vector of the matrix is the constant unit vector, and only the remaining six elements may vary. These six elements are usually represented by $[a, b, c, d, e, f]$. Here is how they are positioned in the matrix:

$$\begin{bmatrix} a & b & 0 \\ c & d & 0 \\ e & f & 1 \end{bmatrix}$$

Next we do a down shift by 100 pixels: $f = 100$

Classes

Matrix

Matrix is a row-major 3x3 matrix used for representing transformations of coordinates throughout MuPDF.

Since all points or pixels reside in a two-dimensional space, one column vector of the matrix is the constant unit vector, and only the remaining six elements may vary. These six elements are usually represented by $[a, b, c, d, e, f]$. Here is how they are positioned in the matrix:

$$\begin{bmatrix} a & b & 0 \\ c & d & 0 \\ e & f & 1 \end{bmatrix}$$

Flipping

Flip the page left-right ($a = -1$)

Classes

Matrix

Matrix is a row-major 3x3 matrix used for representing transformations of coordinates throughout MuPDF. Since all points or pixels reside in a two-dimensional space, one column vector of the matrix is the constant unit vector, and only the remaining six elements may vary. These six elements are usually represented by $[a, b, c, d, e, f]$. Here is how they are positioned in the matrix:

$$\begin{bmatrix} a & b & 0 \\ c & d & 0 \\ e & f & 1 \end{bmatrix}$$

Flip up-down ($d = -1$)

Classes

Matrix

Matrix is a row-major 3x3 matrix used for representing transformations of coordinates throughout MuPDF. Since all points or pixels reside in a two-dimensional space, one column vector of the matrix is the constant unit vector, and only the remaining six elements may vary. These six elements are usually represented by $[a, b, c, d, e, f]$. Here is how they are positioned in the matrix:

$$\begin{bmatrix} e & f & 1 \\ c & d & 0 \\ a & b & 0 \end{bmatrix}$$

Shearing

First a shear in Y direction ($b = 0.5$)

Classes

Matrix

Matrix is a row-major 3x3 matrix used for representing transformations of coordinates throughout MuPDF. With matrices you can manipulate the rendered image of a page in a variety of ways: (parts of) pages can be rotated, zoomed, flipped, sheared and shifted by setting some or all of just six numerical values. Since all points or pixels live in a two-dimensional space, one column vector of that matrix is a constant unit vector, and only the remaining six elements are used for manipulations. These six elements are usually represented by $[a, b, c, d, e, f]$. Here is how they are positioned in the matrix:

$$\begin{bmatrix} a & b & 0 \\ c & d & 0 \\ e & f & 1 \end{bmatrix}$$

It should be noted, that

- the below methods are just convenience functions. Even manipulating $[a, b, c, d, e, f]$
- all manipulations can be combined - you can even

Methods

Matrix `in`

Matrix `in`

Second a shear in X direction ($c = 0.5$)

Matrix

Matrix is a row-major 3x3 matrix used for image transformations in MuPDF. With matrices you can manipulate the rendered image of a page in a variety of ways: (parts of) pages can be rotated, zoomed, flipped, sheared and shifted by setting some or all of just six numerical values.


Since all points or pixels live in a two-dimensional space, one column vector of that matrix is a constant unit vector, and only the remaining six elements are used for manipulations. These six elements are usually represented by $[a, b, c, d, e, f]$. Here is how they are positioned in the matrix:

$$\begin{bmatrix} a & b & 0 \\ c & d & 0 \\ e & f & 1 \end{bmatrix}$$

It should be noted, that

Rotating

Finally a rotation by 60 degrees (`preRotate(60)`)



Outline

`outline` is a property of `Document`. If not `None`, it stands for the first outline item of the document. Its properties in turn define the characteristics of this item and also point to other outline items in "horizontal" direction by property `Outline.next` to the next item of same level, and "downwards" by property `Outline.down` to the next item one level lower. The full tree of all outline items for e.g. a conventional table of contents can be recovered by following these "pointers".

Method / Attribute	Short Description
<code>Outline.down</code>	Next item downwards
<code>Outline.next</code>	Next item same level
<code>Outline.dest</code>	Link destination
<code>Outline.title</code>	Title
<code>Outline.saveText()</code>	Prints a conventional table of contents to a file
<code>Outline.saveXML()</code>	Prints an XML-like table of contents to a file

Class API

class `Outline`

down

The next outline item on the next level down. Is `None` if the item has no children.

Return type: `Outline`

next

The next outline item at the same level as this item. Is `None` if the item is the last one in its level.

Return type: `Outline`

dest

The destination this entry points to. Can be a place in this or another document, or an internet resource. It can include actions to perform like opening a new window, invoking a javascript or opening another document.

Return type: `linkDest`

title

The item's title as a string or `None`.

Return type: `string`

saveText ()

The chain of outline items is being processed and printed to the file `filename` as a conventional table of contents. Each line of this file has the format `<tab>...<tab><title><tab><page#>`, where the number of leading tabs is (n-1), with n equal to the outline level of the entry. Page numbers are 1-based in this case, while `page# = 0` if and only if the outline entry points to a place outside this document. If no title was specified for this outline entry, it appears as a tab character in this file.

Parameters: `filename (string)` -- Name of the file to write to.

saveXML ()

The chain of outline items is being processed and printed to a file `filename` as an XML-like table of contents. Each line of this file has the format `<outline title="..." page="n"/>`, if the entry has no children. Otherwise the format is `<outline title="..." page="n">`, and child entries will follow. The parent entry will be finished by a line containing `</outline>`.

Parameters: `filename (string)` -- Name of the file to write to.

Page

Page interface, created by `Document.loadPage()`.

Method / Attribute	Short Description
<code>Page.bound()</code>	the page's rectangle
<code>Page.loadLinks()</code>	get the first link of a page
<code>Page.getLinks()</code>	get all links of a page
<code>Page.getText()</code>	extract the text of a page
<code>Page.getPixmap()</code>	create a <code>Pixmap</code> for the page
<code>Page.searchFor()</code>	search for a string on a page
<code>Page.run()</code>	run a page through a device
<code>Page.number</code>	page number

Class API

class Page

bound ()

Determine the a page's rectangle (before transformation).

Return type: `Rect`

loadLinks ()

Get the first link of a page.

Return type: `Link`

Returns: A `Link` or `None` if the page has no links.

getLinks ()

Retrieves **all** links of a page.

Return type: list

Returns: A list of dictionaries or `[]`. Please note that the links are in the order as specified during PDF generation, not necessarily in any "natural" reading order.

getText (output = 'text')

Retrieves the text of a page. Depending on the output parameter, the results of the `TextPage` extract methods are returned. The required objects `TextSheet`, `DisplayList`, `Device` and `TextPage` are automatically created or reused if present.

If `output = 'text'` is specified, the text is returned in the order as specified during PDF creation (which is not necessarily the normal reading order). As this may not always look like expected, consider using the example program `PDF2TextJS.py`. It is based on `extractJSON()` output and re-arranges text according to the Western reading layout convention "from top-left to bottom-right".

In an earlier version, we provided a different taste of `TextPage.extractText()` (parameter `basic=False`), which did a similar job. However, the example program mentioned above is more than 50 times faster, so we deleted this option.

Parameters: **output** (*string*) -- A string indicating the requested text format, one of `text` (default), `html`, `json`, or `xml`.

Return type: string

Returns: The page's text as one string.

getPixmap (matrix = fitz.Identity, colorspace = "RGB")

Creates a `Pixmap` from a given page. The method will automatically create or reuse a `DisplayList` for the page and also create the required drawing `Device`.

Parameters:

- **matrix** (*Matrix*) -- A *Matrix* object. Default is the *Identity* matrix.
- **colorspace** (*string*) -- Defines the required colorspace, one of GRAY, CMYK or RGB (default).

Return type: *Pixmap*

Returns: Pixmap of the page.

searchFor (*text*, *hit_max* = 16)

Searches for *text* on a page. Identical to *TextPage.search()*.

Return type: *list*

Returns: A list of *Rect* rectangles each of which surrounds one occurrence of *text*.

run (*dev*, *transform*)

Run a page through a device.

Parameters:

- **dev** (*Device*) -- Device, obtained from one of the *Device* constructors.
- **transform** (*Matrix*) -- Transformation to apply to the page. May include for example scaling and rotation, see *Matrix.preScale()* and *Matrix.preRotate()*. Set it to *Identity* if no transformation is desired.

number

The page number

Return type: *int*

Pixmap

Pixmaps ("pixel maps") are objects at the heart of MuPDF's rendering capabilities. They represent plane rectangular sets of pixels. Each pixel is described by a number of bytes ("components") plus an alpha (transparency) byte.

In PyMuPDF, there exist several ways to create a pixmap:

- create an empty pixmap based on [Colorspace](#) and [IRect](#) information which can then be filled with e.g. the image of a PDF page,
- create a pixmap from a memory area that contains data in PNG format (e.g. obtained from a file),
- create a pixmap from a memory area that contains plain pixels,
- directly create a pixmap from a PDF page via methods `Page.getPixmap()` or `Document.getPagePixmap()`.

Please have a look at the **example** section to see some pixmap usage "at work".

Method / Attribute	Short Description
<code>Pixmap.clearWith()</code>	clears (parts of) a pixmap
<code>Pixmap.tintWith()</code>	tints a pixmap with a color
<code>Pixmap.gammaWith()</code>	applies a gamma factor to the pixmap
<code>Pixmap.writePNG()</code>	saves a pixmap as a PNG file
<code>Pixmap.writeImage()</code>	saves a pixmap in a variety of image formats
<code>Pixmap.copyPixmap()</code>	copy parts of another pixmap
<code>Pixmap.getSize()</code>	returns the pixmap's total length
<code>Pixmap.getColorSpace()</code>	returns the Colorspace used
<code>Pixmap.getIRect()</code>	returns the IRect used
<code>Pixmap.invertIRect()</code>	invert the pixels of a given area
<code>Pixmap.samples</code>	the components data for all pixels
<code>Pixmap.height</code>	height of the region in pixels
<code>Pixmap.width</code>	width of the region in pixels
<code>Pixmap.x</code>	X-coordinate of top-left corner of pixmap
<code>Pixmap.y</code>	Y-coordinate of top-left corner of pixmap
<code>Pixmap.n</code>	number of bytes per pixel
<code>Pixmap.xres</code>	resolution in X-direction
<code>Pixmap.yres</code>	resolution in Y-direction
<code>Pixmap.interpolate</code>	interpolation method indicator

Class API

class Pixmap

`__init__ (self, colorspace, irect)`

This constructor creates an empty pixmap of a size and an origin specified by the irect object. So for a `fitz.IRect(x0, y0, x1, y1)`, `fitz.Point(x0, y0)` will designate the top left corner of the pixmap.

Parameters:

- **colorspace** ([Colorspace](#)) -- The required colorspace of the pixmap.
- **irect** ([IRect](#)) -- Specifies the pixmap's area and its location.

`__init__ (self, data, len)`

This constructor creates a (non-empty) pixmap from data, which is assumed to contain a PNG image of len bytes length. See example below.

Parameters:

- **data** (*string*) -- Data containing a complete, valid image in PNG format. E.g. this may have been obtained from a statement like `data = open('somepic.png', 'rb').read()`. The origin of the resulting pixmap is (0,0).
- **len** (*int*) -- An integer specifying the length of data.

__init__ (self, colorspace, width, height, samples)

This constructor creates a (non-empty) pixmap from `samples`, which is assumed to contain an image in "plain pixel" format. This means that each pixel is represented by `n` bytes (as controlled by the `colorspace` parameter). The origin of the resulting pixmap is (0,0). This method may be useful to create a full copy of a given pixmap, or when raw image data are being provided by some other program - see examples below.

Parameters:

- **colorspace** (*Colorspace*) -- Colorspace of the image. This crucial parameter controls the interpretation of the `samples` area: for `CS_GRAY`, `CS_RGB` and `CS_CMYK`, 2, 4 or 5 bytes in samples will be assumed to define one pixel, respectively.
- **width** (*int*) -- Width of the image
- **height** (*int*) -- Height of the image
- **samples** (*string*) -- A string containing consecutive bytes describing all pixels of the image.

clearWith (value[, iRect])

Clears an area specified by the `IRect` `iRect` within a pixmap. To clear the whole pixmap omit `iRect`.

Parameters:

- **value** (*int*) -- Values from 0 to 255 are valid. Each color byte of each pixel will be set to this value, while alpha will always be set to 255 (non-transparent). Default is 0 (black).
- **iRect** (*IRect*) -- An `IRect` object specifying the area to be cleared.

tintWith (red, green, blue)

Colorizes (tints) a pixmap with a color provided as a value triple (red, green, blue). Use this method only for `CS_GRAY` or `CS_RGB` colorspace. A `TypeError` exception will otherwise be raised.

If the colorspace is `CS_GRAY`, $(red + green + blue) / 3$ will be taken as the tinting value.

Parameters:

- **red** (*int*) -- The `red` component. Values from 0 to 255 are valid.
- **green** (*int*) -- The `green` component. Values from 0 to 255 are valid.
- **blue** (*int*) -- The `blue` component. Values from 0 to 255 are valid.

gammaWith (gamma)

Applies a gamma factor to a pixmap, i.e. lightens or darkens it.

Parameters: **gamma** (*float*) -- `gamma = 1.0` does nothing, `gamma < 1.0` lightens, `gamma > 1.0` darkens the image.

invertIRect (iRect)

Invert the color of all pixels in an area specified by `IRect` `iRect`. To invert everything, omit this parameter.

Parameters: **iRect** (*IRect*) -- The `IRect` to be inverted.

copyPixmap (source, iRect)

Copies the `IRect` part of the `source` pixmap into the corresponding area of this one. The two pixmaps may have different dimensions and different colorspace (provided each is either `CS_GRAY` or `CS_RGB`). The copy mechanism automatically adjusts to any discrepancies between source and target pixmap like so:

If copying from `CS_GRAY` to `CS_RGB`, the source gray-shade value will be put into each of the three rgb component bytes. If the other way round, $(r + g + b) / 3$ will be taken as the gray-shade value of the target.

Between the specified `irect` and the target pixmap's `IRect`, an "intersection" rectangle is first being calculated. Then the corresponding data of this intersection are being copied. If the intersection is empty, nothing will happen.

If you want your `source` pixmap image to land at a specific position of the target, modify its `x` and `y` attributes accordingly before copying. See the example below for how this works.

Parameters:

- **source** (`Pixmap`) -- The pixmap from where to copy.
- **irect** (`IRect`) -- An `IRect` object specifying the area to be copied.

getSize ()

Returns the total length of the pixmap. This will generally equal `len(pix.samples) + 52`. The following will evaluate to `True`: `len(pixmap) == pixmap.getSize()`.

Return type: `int`

getColorspace ()

Returns the colorspace used for constructing the pixmap. Usefull when the pixmap has been created from an image area.

Return type: `Colorspace`

getIRect ()

Returns the `IRect` used for constructing the pixmap. Usefull when the pixmap has been created from an image area.

Return type: `IRect`

writePNG (filename, savealpha=False)

Saves a pixmap as a PNG file. Please note that only grayscale and RGB colorspace can be saved in PNG format (this is not a PyMuPDF restriction).

Parameters:

- **filename** (`string`) -- The filename to save as (the extension `png` must be specified).
- **savealpha** (`bool`) -- Save the alpha channel (`True`) or not (`False` - the default).

writeImage (filename, output="png", savealpha=False)

Saves a pixmap as an image file. This method is an extension to `writePNG()`. Depending on the output chosen, some or all colorspace are supported and different file extensions can be chosen. Please see the table below.

Parameters:

- **filename** (`string`) -- The filename to save to. Depending on the chosen output format, possible file extensions are `.pam`, `.pbm`, `.pgm`, `ppm`, `.pnm`, `.png` and `.tga`.
- **output** (`string`) -- The requested image format. The default is `png` for which this function is equivalent to `writePNG()`. Other possible values are `pam`, `pnm` and `tga`.
- **savealpha** (`bool`) -- Save the alpha channel (`True`) or not (`False` - the default).

samples

The color and transparency values for all pixels. `samples` is a memory area of size `width * height * n` bytes. Each `n` bytes define one pixel. Each successive `n` bytes yield another pixel in scanline order. Subsequent scanlines follow each other with no padding. E.g. for an RGBA colorspace (i.e. `n = 4`) this means, `samples` is a bytearray like `..., R, G, B, A, ...`, and the four byte values `R, G, B, A` define one pixel.

This area can also be used by other graphics libraries like PIL (Python Imaging Library) to do additional processing like saving the pixmap in additional image formats. See example 3.

Return type: `bytearray`

width

The width of the region in pixels. For compatibility reasons, `w` is also supported.

Return type: int

height

The height of the region in pixels. For compatibility reasons, `h` is also supported.

Return type: int

x

X-coordinate of top-left corner

Return type: int

y

Y-coordinate of top-left corner

Return type: int

n

Number of components per pixel. This number depends on (and identifies) the chosen colorspace: `CS_GRAY` = 2, `CS_RGB` = 4, `CS_CMYK` = 5.

Return type: int

xres

Horizontal resolution in dpi (dots per inch).

Return type: int

yres

Vertical resolution in dpi.

Return type: int

interpolate

An information-only boolean flag set to `True` if the image will be drawn using "linear interpolation". If `False` "nearest neighbour sampling" will be used.

Return type: bool

Details on Saving Images with `writeImage()`

The following table shows possible combinations of file extensions, output formats and colorspace of method `writeImage()`.

output =	CS_GRAY	CS_RGB	CS_CMYK
"pam"	.pam	.pam	.pam
"pnm"	.pnm, .pgm	.pnm, .ppm	invalid
"png"	.png	.png	invalid
"tga"	.tga	.tga	invalid

Pixmap Example Code Snippets

Example 1

This shows how pixmaps can be used for purely graphical, non-PDF purposes. The script reads a PNG picture and creates a new PNG file which consist of 3 * 4 tiles of the original one:

```
import fitz
# read in picture image and create a pixmap of it
```

```

pic = open("editra.png", "rb").read()
pix0 = fitz.Pixmap(pic, len(pic))

# calculate target pixmap dimensions and create it
tar_width  = pix0.width * 3          # 3 tiles per row
tar_height = pix0.height * 4         # 4 tiles per column
tarirect   = fitz.IRect(0, 0, tar_width, tar_height)
tar_pix    = fitz.Pixmap(fitz.Colorspace(fitz.CS_RGB), tarirect)
tar_pix.clearWith(90)                # clear target with a lively gray :-


# now fill target with 3 * 4 tiles of input picture
for i in list(range(4)):
    pix0.y = i * pix0.height          # modify input's y coord
    for j in list(range(3)):
        pix0.x = j * pix0.width        # modify input's x coord
        tar_pix.copyPixmap(pix0, pix0.getIRect()) # copy input to new loc
        # save intermediate image to show what is happening
        fn = "target-" + str(i) + str(j) + ".png"
        tar_pix.writePNG(fn)

```

This is the input picture editra.png (taken from the wxPython directory /tools/Editra/pixmaps):



Here is the output, showing some intermediate picture and the final result:



Example 2

This shows how to create a PNG file from a numpy array (several times faster than most other methods):

```

import numpy as np
import fitz
#=====
# create a fun-colored width * height PNG with fitz and numpy
#=====
height = 150
width  = 100
bild=np.ndarray((height, width, 4), dtype=np.uint8)

for i in range(height):

```

```

for j in range(width):
    # one pixel (some fun coloring)
    bild[i, j] = [(i+j)%256, i%256, j%256, 255]

samples = bild.tostring()          # get plain pixel data from numpy array
pix=fitz.Pixmap(fitz.Colorspace(fitz.CS_RGB), width, height, samples)
pix.writePNG("test.png")

```

Example 3

This shows how to interface with PIL / Pillow (the Python Imaging Library), thereby extending the reach of image files that can be processed:

```

import fitz
from PIL import Image

pix = fitz.Pixmap(fitz.csRGB, ...)
... # any code here
# create and save a PIL image
img = Image.frombytes("RGBA", [pix.width, pix.height], str(pix.samples))
img.save(filename, 'jpeg')

# an example for the opposite direction
# create a pixmap from any PIL-supported image file "some_image.xxx"

img = Image.open("some_image.xxx").convert("RGBA")
samples = img.tobytes()
pix = fitz.Pixmap(fitz.csRGB, img.size[0], img.size[1], samples)

```

Point

`Point` represents a point in the plane, defined by its x and y coordinates.

Attribute	Short Description
<code>Point.x</code>	The X-coordinate
<code>Point.y</code>	The Y-coordinate

Class API

`class Point`

`__init__ (self, x, y)`

Constructor, without parameters defaulting to `Point(0.0, 0.0)` ("top left"). Also see the example below.

Parameters:

- `x` (*float*) -- X coordinate of the point
- `y` (*float*) -- Y coordinate of the point

Example:

```
Python 2.7.9 (default, Dec 10 2014, 12:24:55) [MSC v.1500 32 bit (Intel)] on win32
Type "copyright", "credits" or "license()" for more information.
>>> import fitz
>>> point = fitz.Point(25, 30)
>>> point
fitz.Point(25.0, 30.0)
>>>
```

Rect

`Rect` represents a rectangle defined by its top left and its bottom right `Point` objects, in coordinates: $((x_0, y_0), (x_1, y_1))$. Respectively, a rectangle can be defined in one of the two ways: either as a pair of `Point` objects, or as a tuple of four coordinates.

Rectangle borders are always in parallel with the respective X- and Y-axes. A rectangle is called "finite" if $x_0 \leq x_1$ and $y_0 \leq y_1$ is true, else "infinite".

Methods / Attributes	Short Description
<code>Rect.round()</code>	creates the smallest <code>IRect</code> containing <code>Rect</code>
<code>Rect.transform()</code>	transform <code>Rect</code> with a <code>Matrix</code>
<code>Rect.height</code>	<code>Rect</code> height
<code>Rect.width</code>	<code>Rect</code> width
<code>Rect.x0</code>	Top left corner's X-coordinate
<code>Rect.y0</code>	Top left corner's Y-coordinate
<code>Rect.x1</code>	Bottom right corner's X-coordinate
<code>Rect.y1</code>	Bottom right corner's Y-coordinate

Class API

`class Rect`

`__init__ (self, x0, y0, x1, y1)`

Constructor. Without parameters will create the empty rectangle `Rect(0.0, 0.0, 0.0, 0.0)`.

`__init__ (self, p1, p2)`

Constructor. This will create a rectangle from two `Point` objects: `p1` will be taken as the top-left and `p2` as the bottom-right point of the rectangle. See the example below.

`round ()`

Creates the smallest `IRect` that contains `Rect`. This is **not** simply the same as rounding each of the rectangle's coordinates!. Look at the example below.

Return type: `IRect`

`transform (m)`

Transforms `Rect` with a `Matrix`.

Parameters: `m` -- A `Matrix` to be used for the transformation.

Return type: `Rect`

`width`

Contains the width of the rectangle. Equals $x_1 - x_0$.

Return type: float

`height`

Contains the height of the rectangle. Equals $y_1 - y_0$.

Return type: float

`x0`

X-coordinate of the top left corner.

Type: float

`y0`

Y-coordinate of the top left corner.

Type: float

x1

X-coordinate of the bottom right corner.

Type: float

y1

Y-coordinate of the bottom right corner.

Type: float

Example 1:

```
Python 2.7.9 (default, Dec 10 2014, 12:24:55) [MSC v.1500 32 bit (Intel)] on win32
Type "copyright", "credits" or "license()" for more information.
>>> import fitz
>>> p1 = fitz.Point(10, 10)
>>> p2 = fitz.Point(300, 450)
>>> rect1 = fitz.Rect(p1, p2)
>>> rect1
fitz.Rect(10.0, 10.0, 300.0, 450.0)
>>> rect2 = fitz.Rect(0, 0, 250, 380)
>>> rect2
fitz.Rect(0.0, 0.0, 250.0, 380.0)
>>>
```

Example 2:

```
Python 2.7.9 (default, Dec 10 2014, 12:24:55) [MSC v.1500 32 bit (Intel)] on win32
Type "copyright", "credits" or "license()" for more information.
>>> import fitz
>>> r = fitz.Rect(0.5, -0.01, 123.88, 455.123456)
>>> r
fitz.Rect(0.5, -0.009999999776482582, 123.87999725341797, 455.1234436035156)
>>> r.round()
fitz.IRect(0, -1, 124, 456)
>>>
```

As can be seen, `r.round().x0 = int(math.floor(r.x0))`, `r.round().y0 = int(math.floor(r.y0))`, `r.round().x1 = int(math.ceil(r.x1))` and `r.round().y1 = int(math.ceil(r.y1))`.

TextPage

`TextPage` represents the text of a page.

Method	Short Description
<code>TextPage.extractText()</code>	Extract the page's plain text
<code>TextPage.extractHTML()</code>	Extract the page's text in HTML format
<code>TextPage.extractJSON()</code>	Extract the page's text in JSON format
<code>TextPage.extractXML()</code>	Extract the page's text in XML format
<code>TextPage.search()</code>	Search for a string in the page

Class API

`class TextPage`

`extractText()`

Extract the text from a `TextPage` object. Returns a string of the page's complete text. No attempt is being made to adhere to a natural reading sequence: the text is returned UTF-8 encoded and in the same sequence as the PDF creator specified it. If this looks awkward for your PDF file, consider using program that re-arranges the text according to a more familiar layout, e.g. `PDF2TextJS.py` in the examples directory.

Return type: string

`extractHTML()`

Extract the text from a `TextPage` object in HTML format. This version contains some more formatting information about how the text is being displayed on the page. See the tutorial chapter for an example.

Return type: string

`extractJSON()`

Extract the text from a `TextPage` object in JSON format. This version contains significantly more formatting information about how the text is being displayed on the page. It is almost as complete as the `extractXML` version, except that positioning information is detailed down to the span level, not to a single character. See the tutorial chapter for an example. To process the returned JSON text use one of the json modules like `json`, `simplejson`, `ujson`, `cjson`, etc. See example program `PDF2TextJS.py` for how to do that.

Return type: string

`extractXML()`

Extract the text from a `TextPage` object in XML format. This contains complete formatting information about every single text character on the page: font, size, line, paragraph, location, etc. This may easily reach several hundred kilobytes of uncompressed data for a text oriented page. See the tutorial chapter for an example.

Return type: string

`search(string, hit_max = 16)`

Search for the string `string`.

Parameters:

- **string** (*string*) -- The string to search for.
- **hit_max** (*int*) -- Maximum number of expected hits (default 16).

Return type: list

Returns: A python list. If not empty, each element of the list is a [Rect](#) (without transformation) surrounding a found `string` occurrence.

TextSheet

`TextSheet` contains a list of distinct text styles used on a page (or a series of pages).

Constants and Enumerations

Constants and enumerations of MuPDF as implemented by PyMuPDF. If your import statement was `import fitz` then each of the following variables `var` is accessible as `fitz.var`.

Constants

Constant	Description
<code>CS_RGB</code>	1 - Type of Colorspace is RGBA
<code>CS_GRAY</code>	2 - Type of Colorspace is GRAY
<code>CS_CMYK</code>	3 - Type of Colorspace is CMYK
<code>csRGB</code>	= <code>fitz.Colorspace(fitz.CS_RGB)</code>
<code>csGRAY</code>	= <code>fitz.Colorspace(fitz.CS_GRAY)</code>
<code>csCMYK</code>	= <code>fitz.Colorspace(fitz.CS_CMYK)</code>
<code>VersionBind</code>	'1.8.0' - string: version of PyMuPDF (this binding)
<code>VersionFitz</code>	'1.8' - string: version of MuPDF

Enumerations

Possible values of `linkDest.kind` (link destination type). For details consult [Adobe PDF Reference sixth edition 1.7 November 2006](#), chapter 8.2 on page 581 ff.

Value	Description
<code>LINK_NONE</code>	0 - No destination
<code>LINK_GOTO</code>	1 - Points to a place in this document
<code>LINK_URI</code>	2 - Points to an URI
<code>LINK_LAUNCH</code>	3 - Launch (open) another document
<code>LINK_NAMED</code>	4 - Perform some action
<code>LINK_GOTOR</code>	5 - Points to another document

Possible values of `linkDest.flags` (link destination flags). **Attention:** The rightmost byte of this integer is a bit field. The values represent boolean indicators showing whether the associated statement is `True`.

Value	Description
<code>LINK_FLAG_L_VALID</code>	1 (bit 0) Top left x value is valid
<code>LINK_FLAG_T_VALID</code>	2 (bit 1) Top left y value is valid

LINK_FLAG_R_VALID	4 (bit 2) Bottom right x value is valid
LINK_FLAG_B_VALID	8 (bit 3) Bottom right y value is valid
LINK_FLAG_FIT_H	16 (bit 4) Horizontal fit
LINK_FLAG_FIT_V	32 (bit 5) Vertical fit
LINK_FLAG_R_IS_ZOOM	64 (bit 6) Bottom right x is a zoom figure

Appendix 1: Performance

We have tried to get an impression on PyMuPDF's performance. While we know this is very hard and a fair comparison is almost impossible, we feel that we at least should provide some quantitative information to justify our bold comments on MuPDF's **top performance**.

Following are three sections that deal with different aspects of performance:

- document parsing
- text extraction
- image rendering

In each section, the same fixed set of PDF files is being processed by a set of tools. The set of tools varies - for reasons we will explain in the section.

Here is the list of files we are using. Each file name is accompanied by further information: **size** in bytes, number of **pages**, number of bookmarks (**toc** entries), number of **links**, **text** size as a percentage of file size, **KB** per page, PDF **version** and remarks. **text %** and **KB index** are indicators for whether a file is text or graphics oriented: e.g. `Adobe.pdf` and `PyMuPDF.pdf` are clearly text oriented, all other files contain many more images.

name	size	pages	toc size	links	text %	KB index	version	remarks
Adobe.pdf	32.472.771	1.310	794	32.096	8,0%	24	PDF 1.6	linearized, text oriented, many links / bookmarks
Evolution.pdf	13.497.490	75	15	118	1,1%	176	PDF 1.4	graphics oriented
PyMuPDF.pdf	479.011	47	60	491	13,2%	10	PDF 1.4	text oriented, many links
sdw_2015_01.pdf	14.668.972	100	36	0	2,5%	143	PDF 1.3	graphics oriented
sdw_2015_02.pdf	13.295.864	100	38	0	2,7%	130	PDF 1.4	graphics oriented
sdw_2015_03.pdf	21.224.417	108	35	0	1,9%	192	PDF 1.4	graphics oriented
sdw_2015_04.pdf	15.242.911	108	37	0	2,7%	138	PDF 1.3	graphics oriented
sdw_2015_05.pdf	16.495.887	108	43	0	2,4%	149	PDF 1.4	graphics oriented
sdw_2015_06.pdf	23.447.046	100	38	0	1,6%	229	PDF 1.4	graphics oriented
sdw_2015_07.pdf	14.106.982	100	38	2	2,6%	138	PDF 1.4	graphics oriented
sdw_2015_08.pdf	12.321.995	100	37	0	3,0%	120	PDF 1.4	graphics oriented
sdw_2015_09.pdf	23.409.625	100	37	0	1,5%	229	PDF 1.4	graphics oriented
sdw_2015_10.pdf	18.706.394	100	24	0	2,0%	183	PDF 1.5	graphics oriented
sdw_2015_11.pdf	25.624.266	100	20	0	1,5%	250	PDF 1.4	graphics oriented
sdw_2015_12.pdf	19.111.666	108	36	0	2,1%	173	PDF 1.4	graphics oriented

Decimal point and comma follow European convention

Part 1: Parsing

How fast is a PDF file read and its content parsed for further processing? The sheer parsing performance cannot directly be compared, because batch utilities always execute a requested task in one go, front to end. `pdfwrw` too, has a `lazy` strategy for parsing, meaning it only parses those parts of a document that are required in any moment.

We therefore measure the time to copy a PDF file to an output file, and doing nothing else.

These were the tools

All tools are either platform independant, or at least can run on Windows and Unix / Linux (`pdftk`).

Poppler is missing here, because it specifically is a Linux tool set, although we know there exist Windows ports (created with considerable effort apparently). Technically, it is a C/C++ library, for which a Python binding exists - in so far it is somewhat comparable to PyMuPDF. But Poppler in contrast is tightly coupled to **Qt** and **Cairo**. We may still include it in future, when a more handy Windows installation is available. We have seen however some [analysis](#), that hints at a much lower performance than MuPDF. Our comparison of text extraction speeds also show a much lower performance of Poppler's PDF code base **Xpdf**.

Image rendering of MuPDF also is about three times faster than the one of Xpdf when comparing the command line tools `mudraw` of MuPDF and `pdftopng` of Xpdf - see part 3 of this chapter.

Tool	Description
PyMuPDF	tool of this manual, appearing as "fitz" in reports
pdfwrw	a pure Python tool, can be used as frontend to ReportLab and rst2pdf

PyPDF2	a pure Python tool with a very complete function set
pdftk	a command line utility with numerous functions

This is how each of the tools is being used with the test:

PyMuPDF:

```
doc = fitz.Document("input.pdf")
doc.save("output.pdf")
```

pdfrw:

```
doc = PdfReader("input.pdf")
writer = PdfWriter()
writer.trailer = doc
writer.write("output.pdf")
```

PyPDF2:

```
pdfmerge = PyPDF2.PdfFileMerger()
pdfmerge.append("input.pdf")
pdfmerge.write("output.pdf")
pdfmerge.close()
```

pdftk:

```
pdftk input.pdf output output.pdf
```

Observations

These are our run time findings (in **seconds**, please note the European number convention: meaning of decimal point and comma is reversed):

Runtimes	tool			
	fitz	pdfrw	pdftk	PyPDF2
Adobe.pdf	8,58	20,27	109,22	672,13
Evolution.pdf	0,14	0,41	1,20	1,32
PyMuPDF.pdf	0,20	0,17	0,74	0,88
sdw_2015_01.pdf	0,34	1,13	5,37	5,14
sdw_2015_02.pdf	0,40	1,49	6,48	6,49
sdw_2015_03.pdf	0,92	2,62	11,19	11,61
sdw_2015_04.pdf	0,45	2,14	7,16	6,91
sdw_2015_05.pdf	0,42	1,64	7,37	7,34
sdw_2015_06.pdf	1,24	3,22	13,67	13,86
sdw_2015_07.pdf	0,92	2,06	9,33	9,25
sdw_2015_08.pdf	0,50	1,86	8,31	8,31
sdw_2015_09.pdf	1,14	2,31	9,86	9,99
sdw_2015_10.pdf	0,86	1,80	3,10	6,38
sdw_2015_11.pdf	6,73	12,39	35,11	58,49
sdw_2015_12.pdf	0,59	2,10	8,51	8,91
total times	23,43	55,61	236,62	827,00

Timing Ratios			
1,00	2,37	10,10	35,30
	1,00	4,25	14,87
		1,00	3,50

If we leave out the Adobe manual, this table looks like

Runtimes	tool			
file	fitz	pdfrw	pdftk	PyPDF2
Evolution.pdf	0,14	0,41	1,20	1,32
PyMuPDF.pdf	0,20	0,17	0,74	0,88
sdw_2015_01.pdf	0,34	1,13	5,37	5,14
sdw_2015_02.pdf	0,40	1,49	6,48	6,49
sdw_2015_03.pdf	0,92	2,62	11,19	11,61
sdw_2015_04.pdf	0,45	2,14	7,16	6,91
sdw_2015_05.pdf	0,42	1,64	7,37	7,34
sdw_2015_06.pdf	1,24	3,22	13,67	13,86
sdw_2015_07.pdf	0,92	2,06	9,33	9,25
sdw_2015_08.pdf	0,50	1,86	8,31	8,31
sdw_2015_09.pdf	1,14	2,31	9,86	9,99
sdw_2015_10.pdf	0,86	1,80	3,10	6,38
sdw_2015_11.pdf	6,73	12,39	35,11	58,49
sdw_2015_12.pdf	0,59	2,10	8,51	8,91
total times	14,85	35,34	127,40	154,87

Timing Ratios			
1,00	2,38	8,58	10,43
	1,00	3,60	4,38
		1,00	1,22

PyMuPDF is by far the fastest: on average 2.4 times faster than the second best (the pure Python tool pdfrw, **chapeau pdfrw!**), and 10 times faster than the command line tool pdftk.

Where PyMuPDF only requires less than 24 seconds to process all files, pdftk affords itself almost 4 minutes.

By far the slowest tool is PyPDF2 - it is more than 35 times slower than PyMuPDF and 15 times slower than pdfrw! The main reason for PyPDF2's bad look comes from the Adobe manual. It obviously is slowed down by the linear file structure and the immense amount of bookmarks of this file. If we take out this special case, then PyPDF2 is 10.5 times slower than PyMuPDF, 4.4 times slower than pdfrw and 1.2 times slower than pdftk.

If we look at the output PDFs, there is one surprise:

Each tool created a PDF of similar size as the original. Apart from the Adobe case, PyMuPDF always created the smallest output.

Adobe's manual is an exception: The pure Python tools **reduced** its size by more than 20% (and yielded a document which is no longer linearized)!

PyMuPDF and pdftk in contrast **drastically increased** the size by 40% to about 50 MB (also no longer linearized).

So far, we have no explanation of what is happening here.

Part 2: Text Extraction

We also have compared text extraction speed with other tools.

The following table shows a run time comparison. PyMuPDF's methods appear as "fitz (TEXT)" and "fitz (JSON)" respectively. The tool `pdftotext.exe` of the [Xpdf](#) toolset appears as "xpdf".

- **extractText():** basic text extraction without layout re-arrangement (using `GetText(..., output = "text")`)
- **pdftotext:** a command line tool of the [Xpdf](#) toolset (which also is the basis of [Poppler's library](#))
- **extractJSON():** text extraction with layout information (using `GetText(..., output = "json")`)
- **pdfminer:** a pure Python PDF tool specialized on text extraction tasks

All tools have been used with their most basic, fanciless functionality - no layout re-arrangements, etc.

For demonstration purposes, we have included a version of `GetText(doc, output = "json")`, that also re-arranges the output according to occurrence on the page.

Here are the results using the same test files as above (again: decimal point and comma reversed):

Text Extraction	Tool				
File	1 fitz (TEXT)	2 fitz (bare JSON)	3 fitz (sorted JSON)	4 xpdf	5 pdfminer
Adobe.pdf	3,86	4,94	5,73	10,57	213,10
Evolution.pdf	0,21	0,26	0,31	0,48	12,70
PyMuPDF.pdf	0,08	0,11	0,13	0,96	4,54
sdw_2015_01.pdf	0,74	0,86	0,98	1,29	42,48
sdw_2015_02.pdf	0,81	1,00	1,04	1,39	46,84
sdw_2015_03.pdf	1,47	1,63	1,71	2,34	149,51
sdw_2015_04.pdf	0,94	1,10	1,21	1,67	79,11
sdw_2015_05.pdf	0,80	0,96	1,09	1,32	47,47
sdw_2015_06.pdf	1,50	1,65	1,77	2,22	135,13
sdw_2015_07.pdf	0,79	0,93	1,02	1,40	53,87
sdw_2015_08.pdf	0,78	0,90	1,02	1,33	46,78
sdw_2015_09.pdf	1,50	1,60	1,73	2,29	155,52
sdw_2015_10.pdf	0,90	1,02	1,09	1,86	72,03
sdw_2015_11.pdf	2,03	2,15	2,23	4,31	161,43
sdw_2015_12.pdf	0,88	1,02	1,12	1,61	68,03
Total Times	17,28	20,11	22,17	35,04	1288,52

Timing Relationships				
1,0	1,2	1,3	2,0	74,6
	1,0	1,1	1,7	64,1
		1,0	1,6	58,1
			1,0	36,8

Again, (Py-) MuPDF is the fastest around. It is two times faster than xpdf.

JSON output is 1.7 times faster than xpdf, and even the "re-arranging" version is 1.6 times faster.

pdfminer, as a pure Python solution, of course is comparatively slow: MuPDF is 75 (64, 58) times faster and xpdf is 37 times faster. These observations in order of magnitude coincide with the statements on this [web site](#).

Part 3: Image Rendering

We have tested rendering speed of MuPDF against the `pdftopng.exe`, a command line tool of the **Xpdf** toolset, which is the PDF code basis of **Poppler**.

MuPDF invocation using a resolution of 150 pixels (Xpdf default):

```
mutool draw -o t%d.png -r 150 file.pdf
```

Xpdf invocation:

```
pdftopng.exe file.pdf ./
```

The resulting runtimes can be found here (again: meaning of decimal point and comma reversed):

Run Time (sec)	Tool	
File	mudraw	xpdf
Adobe.pdf	107,99	505,27
Evolution.pdf	41,05	108,33
PyMuPDF.pdf	5,74	21,82
sdw_2015_01.pdf	29,93	76,81
sdw_2015_02.pdf	29,82	74,68
sdw_2015_03.pdf	31,95	85,89
sdw_2015_04.pdf	29,16	78,09
sdw_2015_05.pdf	30,63	77,56
sdw_2015_06.pdf	30,77	87,89
sdw_2015_07.pdf	33,06	78,74
sdw_2015_08.pdf	31,26	75,95
sdw_2015_09.pdf	35,48	84,37
sdw_2015_10.pdf	29,47	77,13
sdw_2015_11.pdf	31,97	80,96
sdw_2015_12.pdf	31,30	80,68
Total	529,61	1594,18

MuPDF is between 2.7 and 4.7 (on average 3.0) times faster than Xpdf.

Appendix 2: Details on Text Extraction

This chapter provides background on the text extraction methods of PyMuPDF.

Information of interest are

- what do they provide?
- what do they imply (processing time / data sizes)?

General structure of a *TextPage*

Text information contained in a *TextPage* adheres to the following hierarchy:

```
<page> (width and height)
  <block> (its rectangle)
    <line> (its rectangle)
      <span> (its rectangle and font information)
        <char> (its rectangle, (x, y) coordinates and value)
```

A **text page** consists of blocks (= roughly paragraphs).

A **block** consists of lines.

A **line** consists of spans.

A **span** consists of characters with the same properties. E.g. a different font will cause a new span.

Output of `getText (output="text ")`

This is the plain text output of a page of this tutorial's PDF version:

```
Tutorial

This tutorial will show you the use of MuPDF in Python step by step.

Because MuPDF supports not only PDF, but also XPS, OpenXPS and EPUB formats, so does PyMuPDF.

Nevertheless we will only talk about PDF files for the sake of brevity.
...
```

Output of `getText (output="html ")`

The HTML version looks like this:

```
<div class="page">
<div class="block"><p>
<div class="metaline"><div class="line"><div class="cell" style="width:0%;align:left"><span
</div></p></div>
<div class="block"><p>
<div class="line"><div class="cell" style="width:0%;align:left"><span class="s1">This tutorial
</div></p></div>
<div class="block"><p>
<div class="line"><div class="cell" style="width:0%;align:left"><span class="s1">Because MuP
<div class="line"><div class="cell" style="width:0%;align:left"><span class="s1">Nevertheles
</div></p></div>
...
```

Output of `getText (output="json ")`

JSON output looks like so:

```
{
  "len":35,"width":595.2756,"height":841.8898,
  "blocks":[
```

```

{
  "type": "text",
  "bbox": [40.01575, 53.730354, 98.68775, 76.08236],
  "lines": [
    {
      "bbox": [40.01575, 53.730354, 98.68775, 76.08236],
      "spans": [
        {
          "bbox": [40.01575, 53.730354, 98.68775, 76.08236],
          "text": "Tutorial"
        }
      ]
    }
  ]
},
{
  "type": "text",
  "bbox": [40.01575, 79.300354, 340.6957, 93.04035],
  "lines": [
    {
      "bbox": [40.01575, 79.300354, 340.6957, 93.04035],
      "spans": [
        {
          "bbox": [40.01575, 79.300354, 340.6957, 93.04035],
          "text": "This tutorial will show you the use of MuPDF in Python step by step."
        }
      ]
    }
  ]
},
...

```

Output of `getText (output="xml")`

Now the XML version:

```

<page width="595.2756" height="841.8898">
<block bbox="40.01575 53.730354 98.68775 76.08236">
<line bbox="40.01575 53.730354 98.68775 76.08236">
<span bbox="40.01575 53.730354 98.68775 76.08236" font="Helvetica-Bold" size="16">
<char bbox="40.01575 53.730354 49.79175 76.08236" x="40.01575" y="70.85036" c="T"/>
<char bbox="49.79175 53.730354 59.56775 76.08236" x="49.79175" y="70.85036" c="u"/>
<char bbox="59.56775 53.730354 64.89575 76.08236" x="59.56775" y="70.85036" c="t"/>
<char bbox="64.89575 53.730354 74.67175 76.08236" x="64.89575" y="70.85036" c="o"/>
<char bbox="74.67175 53.730354 80.89575 76.08236" x="74.67175" y="70.85036" c="r"/>
<char bbox="80.89575 53.730354 85.34375 76.08236" x="80.89575" y="70.85036" c="i"/>
<char bbox="85.34375 53.730354 94.23975 76.08236" x="85.34375" y="70.85036" c="a"/>
<char bbox="94.23975 53.730354 98.68775 76.08236" x="94.23975" y="70.85036" c="l"/>
</span>
</line>
</block>
<block bbox="40.01575 79.300354 340.6957 93.04035">
<line bbox="40.01575 79.300354 340.6957 93.04035">
<span bbox="40.01575 79.300354 340.6957 93.04035" font="Helvetica" size="10">
<char bbox="40.01575 79.300354 46.12575 93.04035" x="40.01575" y="90.050354" c="T"/>
<char bbox="46.12575 79.300354 51.685753 93.04035" x="46.12575" y="90.050354" c="h"/>
<char bbox="51.685753 79.300354 53.90575 93.04035" x="51.685753" y="90.050354" c="i"/>
<char bbox="53.90575 79.300354 58.90575 93.04035" x="53.90575" y="90.050354" c="s"/>
<char bbox="58.90575 79.300354 61.685753 93.04035" x="58.90575" y="90.050354" c=" "/>
<char bbox="61.685753 79.300354 64.46575 93.04035" x="61.685753" y="90.050354" c="t"/>
<char bbox="64.46575 79.300354 70.02576 93.04035" x="64.46575" y="90.050354" c="u"/>
<char bbox="70.02576 79.300354 72.805756 93.04035" x="70.02576" y="90.050354" c="t"/>
<char bbox="72.805756 79.300354 78.36575 93.04035" x="72.805756" y="90.050354" c="o"/>
<char bbox="78.36575 79.300354 81.695755 93.04035" x="78.36575" y="90.050354" c="r"/>
<char bbox="81.695755 79.300354 83.91576 93.04035" x="81.695755" y="90.050354" c="i"/>
...

```

Resource Requirements

The four text extraction methods of a [TextPage](#) differ significantly: in terms of information they supply (see above), and in terms of resource requirements. More information of course means that more processing is required and a higher data volume is generated.

For testing performance, we have run several example PDFs through these methods and found the following information. This data is not statistically secured in any way - just take it as an idea for what you should expect to see.

As a low end example we took this manual's PDF version (45+ pages, text oriented, 500 KB). The high end case was Adobe's PDF manual (1310 pages, text oriented, 32 MB). The other test cases were [Spektrum](#) magazines of the year 2015 (the German version of Scientific American, 100+ pages, text with lots of complex interspersed images, 10 to 25 MB each).

Performance

Performance of text extraction has improved significantly in MuPDF 1.8! As of updating this documentation (mid November 2015), data hint at an improvement factor greater than 2. Especially the complex extraction methods now have a much lower effort penalty.

On a higher level Win10 machine (8 processors at 4 GHz, 8 GB RAM), `extractXML()` needs anything between 0.2 and 0.5 seconds per page. This means that you can extract extremely detailed text information of a complex 100-page magazine in less than a minute. This is faster than some other free text extraction tools like e.g. [Nitro 3](#).

With `PDF2TextJS.py` of the example directory, you have a high performance text extraction utility with a high layout faithfulness!

Data Sizes

The sizes of the returned text strings follow this pattern (`extractText()` is set to 1):

(Text : HTML : JSON : XML) ~ (1 : 4 : 6 : 87)

The number 87 for `extractXML()` corresponds to values between 200 and 400 KB per page.

Index

`__init__()` (Colorspace method)
(Device method) [1]
(Document method) [1]
(IRect method)
(Matrix method)
(Pixmap method) [1] [2]
(Point method)
(Rect method) [1]

A

`a` (Matrix attribute)
`authenticate()` (Document method)
`author` (built-in variable)

B

`b` (Matrix attribute)
`bound()` (Page method)

C

`c` (Matrix attribute)
`clearWith()` (Pixmap method)
`close()` (Document method)
`Colorspace` (built-in class)
`copyPixmap()` (Pixmap method)
`creationDate` (built-in variable)
`creator` (built-in variable)
`CS_CMYK` (built-in variable)
`CS_GRAY` (built-in variable)
`CS_RGB` (built-in variable)
`csCMYK` (built-in variable)
`csGRAY` (built-in variable)
`csRGB` (built-in variable)

D

`d` (Matrix attribute)
`dest` (Link attribute)
(Outline attribute)
(linkDest attribute)
`Device` (built-in class)
`DisplayList` (built-in class)

`DisplayList()` (`DisplayList.fitz` method)
`Document` (built-in class)
`down` (Outline attribute)

E

`e` (Matrix attribute)
`encryption` (built-in variable)
`extractHTML()` (TextPage method)
`extractJSON()` (TextPage method)
`extractText()` (TextPage method)
`extractXML()` (TextPage method)

F

`f` (Matrix attribute)
`fileSpec` (linkDest attribute)
`flags` (linkDest attribute)
`format` (built-in variable)

G

`gammaWith()` (Pixmap method)
`getColorspace()` (Pixmap method)
`getIRect()` (Pixmap method)
`getLinks()` (Page method)
`getPagePixmap()` (Document method)
`getPageText()` (Document method)
`getPixmap()` (Page method)
`getRect()` (IRect method)
`getSize()` (Pixmap method)
`getText()` (Page method)
`getToC()` (Document method)

H

`height` (IRect attribute)
(Pixmap attribute)
(Rect attribute)

I

`interpolate` (Pixmap attribute)
`invertIRect()` (Pixmap method)
`IRect` (built-in class)
`isClosed` (Document attribute)
`isEncrypted` (Document attribute)
`isMap` (linkDest attribute)

isUri (linkDest attribute)

K

keywords (built-in variable)

kind (linkDest attribute)

L

Link (built-in class)

LINK_FLAG_B_VALID (built-in variable)

LINK_FLAG_FIT_H (built-in variable)

LINK_FLAG_FIT_V (built-in variable)

LINK_FLAG_L_VALID (built-in variable)

LINK_FLAG_R_IS_ZOOM (built-in variable)

LINK_FLAG_R_VALID (built-in variable)

LINK_FLAG_T_VALID (built-in variable)

LINK_GOTO (built-in variable)

LINK_GOTOR (built-in variable)

LINK_LAUNCH (built-in variable)

LINK_NAMED (built-in variable)

LINK_NONE (built-in variable)

LINK_URI (built-in variable)

linkDest (built-in class)

loadLinks() (Page method)

loadPage() (Document method)

lt (linkDest attribute)

M

Matrix (built-in class)

metadata (Document attribute)

modDate (built-in variable)

N

n (Pixmap attribute)

name (Document attribute)

named (linkDest attribute)

needsPass (Document attribute)

newWindow (linkDest attribute)

next (Link attribute)

(Outline attribute)

number (Page attribute)

O

object (Device attribute)

Outline (built-in class)

outline (Document attribute)

P

Page (built-in class)

page (linkDest attribute)

pageCount (Document attribute)

Pixmap (built-in class)

Point (built-in class)

preRotate() (Matrix method)

preScale() (Matrix method)

preShear() (Matrix method)

producer (built-in variable)

R

rb (linkDest attribute)

Rect (built-in class)

rect (Link attribute)

round() (Rect method)

run() (DisplayList method)

(Page method)

S

samples (Pixmap attribute)

save() (Document method)

saveText() (Outline method)

saveXML() (Outline method)

search() (TextPage method)

searchFor() (Page method)

subject (built-in variable)

T

TextPage (built-in class)

textpage (Device attribute)

textsheet (Device attribute)

tintWith() (Pixmap method)

title (built-in variable)

(Outline attribute)

transform (Rect attribute)

U

uri (linkDest attribute)

V

VersionBind (built-in variable)

VersionFitz (built-in variable)

W

width (IRect attribute)

(Pixmap attribute)

(Rect attribute)

writelnImage() (Pixmap method)

writePNG() (Pixmap method)

X

x (Pixmap attribute)

x0 (IRect attribute)

(Rect attribute)

x1 (IRect attribute)

(Rect attribute)

xres (Pixmap attribute)

Y

y (Pixmap attribute)

y0 (IRect attribute)

(Rect attribute)

y1 (IRect attribute)

(Rect attribute)

yres (Pixmap attribute)